

Technology for X-Informatics

Kmeans and MapReduce Parallelism

June 19 2013

Geoffrey Fox

gcf@indiana.edu

<http://www.infomall.org/>

Associate Dean for Research, School of Informatics
and Computing

Indiana University Bloomington

2013

Big Data Ecosystem in One Sentence

Use **Clouds** running **Data Analytics Collaboratively**
processing **Big Data** to solve problems in
X-Informatics (or e-X)

X = Astronomy, Biology, Biomedicine, Business, Chemistry, Climate,
Crisis, Earth Science, Energy, Environment, Finance, Health,
Intelligence, Lifestyle, Marketing, Medicine, Pathology, Policy, Radar,
Security, Sensor, Social, Sustainability, Wealth and Wellness with
more fields (physics) defined implicitly
Spans Industry and Science (research)

Education: **Data Science** see recent New York Times articles
<http://datascience101.wordpress.com/2013/04/13/new-york-times-data-science-articles/>



Climate Informatics
network

How Wealth Informatics can help
with your financial freedom?



Biomedical Informatics

Computer Applications in Health Care
and Biomedicine

AstroInformatics2012


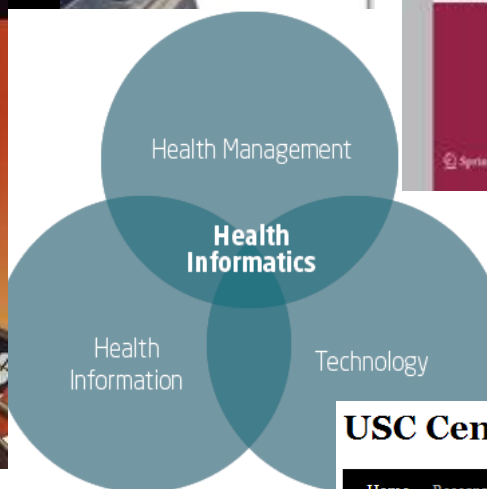

Redmond, WA, September 10 - 14, 2012

Journal of
Pathology
Informatics


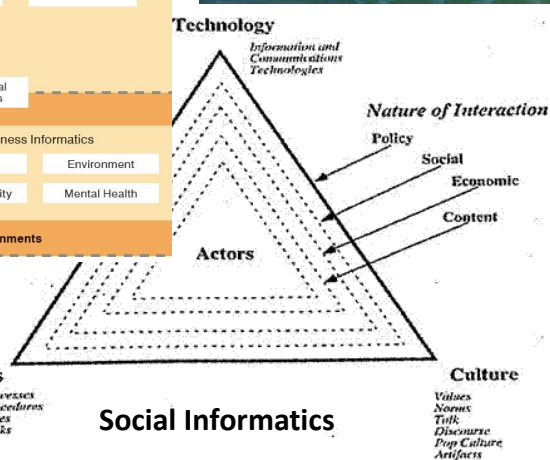
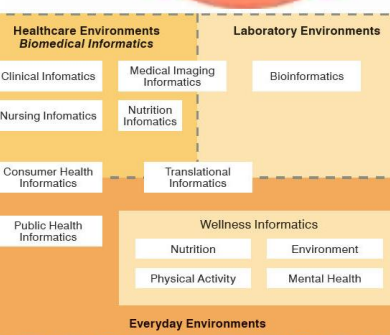
Intelligence and
Security Informatics



RICHARD E. NEAPOLITAN • XIA JIANG



Opportunities and Challenges
in Crisis Informatics



Noella Penelope Greer (Ed.)

Business Informatics
Information technology, Management,

policy informatics network




USC Center For Energy Informatics

Home Research Publications Smart

About the Center


Welcome to the Center For Energy Informatics (CEI) at USC, an Organized Research Unit (ORU) housed in the [Viterbi School of Engineering](#). Energy Informatics is the application of

Lifestyle Informatics



Lifestyle Informatics: Let people

The study Lifestyle Informatics is about... this bachelor including applied psycholog... knowledge about language and informatic... short better. Lifestyle Informatics: let peo... [Lifestyle Informatics](#)



MapReduce Kmeans in Python

Slightly Changed Sequential version

- We change to allow MapReduce implementation
- Overall iteration over starting positions
 - Initialize Centroids
 - Iterate until converged
 - Call **map** to find association and do first step in find centroids as sum over point vectors and distortion
 - **Reduce** step: Divide summed centers by points in center to get centroid
 - Do other book keeping
 - Delete zero size clusters
 - Check convergence
- Return best solution based on Quality criterion

Parallel MapReduce Version

```
lseven = np.empty([tot], dtype=bool)
for i in arange(tot):
    lseven[i] = (i%2 == 0);
obs1 = compress(lseven, obs, 0)
obs2 = compress(logical_not(lseven), obs, 0)
```

```
avg_dist = []
diff = thresh+1.
while diff > thresh:
    #
    if Parallelism == 1:
        code_book, NumPointsinClusters, distortsum, distortmax, NumPoints = Kmeans_map(obs, code_book)
    if Parallelism == 2:
        # Can be Parallel Map Operations
        code_book1, NumPointsinClusters1, distortsum1, distortmax1, NumPoints1 = Kmeans_map(obs1, code_book)
        code_book2, NumPointsinClusters2, distortsum2, distortmax2, NumPoints2 = Kmeans_map(obs2, code_book)
        #
        # Following are 4 Reduction Operations
        # Note maps include local reductions
        code_book = np.add( code_book1, code_book2)
        NumPointsinClusters = np.add( NumPointsinClusters1, NumPointsinClusters2)
        distortsum = distortsum1 + distortsum2
        distortmax = np.maximum(distortmax1, distortmax2)
        NumPoints = NumPoints1 + NumPoints2
        #
        code_book = compress(np.greater(NumPointsinClusters, 0), code_book, 0)
        # remove code_books that didn't have any members
        #
        j = 0
        nc = code_book.shape[0]
        for i in arange(nc):
            if NumPointsinClusters[i] > 0:
                code_book[j,:] = code_book[i,:] / NumPointsinClusters[i]
                j = j + 1
```

Map

Reduce

Map for Kmeans

```
def Kmeans_map(obs, code_book):  
    No = obs.shape[0]  
    nc = code_book.shape[0]  
    # nc is current number of clusters (may decrease if zero clusters last iteration)  
    #  
    #compute membership and distances between obs and code_book  
    obs_code, distort = vq(obs, code_book)  
    distortsum = np.sum(distort)  
    distortmax = np.amax(distort)  
    #  
    # vq returns an indexing array obs_code mapping rows of obs (the points) to code_book (the centroids)  
    # distort is an array of length No that has difference between observation and chosen centroid  
    # vq stands for vector quantization and is provided in SciPy  
    #  
    VectorDimension = obs.shape[1]  
    NewCode_Book = np.zeros([nc, VectorDimension])  
    NumPointsinClusters = np.zeros([nc])  
    for i in arange(nc):  
        #          Loop over clusters labelled with i  
        cell_members = compress(equal(obs_code, i), obs, 0)  
        NumPointsinClusters[i] = cell_members.shape[0]  
        #          Extract Points in this Cluster; extract points whose quantization label is i  
        #  
        NewCode_Book[i] = np.sum(cell_members, 0)  
        #          Calculate centroid of i'th cluster  
    return NewCode_Book, NumPointsinClusters, distortsum, distortmax, No
```

- This routine takes Points in obs and Centroids in code_book and associates each point with nearest centroid
- It calculates Non normalized centroids (sum over vectors in cluster), Number of points in each cluster and convergence measures (mean and max of distance between points and centers)

Comments!

- Extension to P-way parallelism simple
 - Divide data into P parts
 - Run P separate maps with each having all centers
 - Sum (or max) P versions of center sums, Point sums, convergence criterion in Reduce
- Maps can run in parallel
- Reduce would run in a single process except for large levels of parallelism
- Sequential code uses mean directly not sums
- Parallel code does center sums and
- Divides by total number of points in Reduce